

Backprop: How it works, and how it fails

Things to do: Add movies of training etc + movies of GAN/VAE/Pole balancing/RL demos.

James V Stone

Sheffield University

What's the problem?

A deep neural network: Credit assignment

The simplest neural network

Gradient descent: The delta term


Gradient descent with two input units

The Backprop Algorithm

Local minima


Overfitting and generalisation

A Deep Neural Network: Credit Assignment



A deep network with three hidden layers.
This embodies the *credit assignment* or *blame assignment* problem.

The Simplest Neural Network: Definitions



Definitions

x *state* of input unit

w *weight* connecting two units


u total *input* to a unit (e.g. to output unit)

y *state* of output unit

$y = f(u)$ the *activation function* of a unit is f

y desired or *target value* of output unit

The Simplest Neural Network: From input to output



If the input state is x and if the weight of the connection from the input unit to the output unit is w , then the total input u to the output unit is

$$u = wx. \quad (1)$$

In general, the state y of a unit is governed by an *activation function* (i.e. input/output function)

$$y = f(u). \quad (2)$$

The Simplest Neural Network: The Delta Term

Suppose we wish the network to learn to associate an input value of $x = 0.8$ with a target state of $y = 0.2$. Usually, we have no idea of the correct value for the weight w , so we may as well begin by choosing its value at random. Suppose we choose a value $w = 0.4$, so the output state is $y = 0.4 \times 0.8 = 0.32$.

The difference between the output y and the target value y is defined here as the *delta term*:

$$\delta = y - y = 0.32 - 0.2 = 0.12. \quad (3)$$


Ideally, we would like to adjust the weight w so that $\delta = 0$. A standard measure of the error in y is half the squared difference:

$$E = \frac{1}{2} (wx - y)^2 \quad (4)$$

$$= \frac{1}{2} (y - y)^2 \quad (5)$$

$$= \frac{1}{2} \delta^2. \quad (6)$$


The Simplest Neural Network



The optimal weight is $w^* = 0.25$.

The gradient of the error function E at a point w is approximated by $\Delta E / \Delta w$.

The Simplest Neural Network




Given that

$$E = \frac{1}{2} (wx - y)^2, \quad (7)$$

the gradient of the error function at a point w is approximated by $\Delta E / \Delta w$. An exact measure of the gradient is defined by the derivative of E (Equation 7) with respect to w :

$$\frac{dE}{dw} = (wx - y)x \approx \frac{\Delta E}{\Delta w}. \quad (8)$$

The Simplest Neural Network




It will prove useful to write Equation 8 in terms of the delta term:

$$\frac{dE}{dw} = \delta x. \quad (9)$$

The *magnitude* of the gradient indicates the steepness of the slope at w , and the *sign* of the gradient indicates the direction that increases E .

The Simplest Neural Network



The direction of the gradient measured using calculus points uphill, and is called the *direction of steepest ascent*. This means that in order to reduce E , we should change the value of w by a small amount Δw in the *direction of steepest descent*:

$$\Delta w = -\epsilon \frac{dE}{dw} \quad (10)$$

$$= -\epsilon \delta x, \quad (11)$$

where the size of the step is defined by a *learning rate parameter* ϵ .

The Simplest Neural Network: Algorithm

Gradient Descent

Learning One Association

initialise network weight w to random value

set input unit states x to training vector x

set learning to true

while *learning* **do**

 get state of output unit $y = wx$

 get delta term $\delta = y - y$

 get weight gradient for input vector $dE/dw = \delta x$

 get change in weight $\Delta w = -\epsilon dE/dw$

 update weight $w \leftarrow w + \Delta w$

if *gradient* $dE/dw \approx 0$ **then**

 set learning to false

end

end

Gradient Descent for a Network with Two Input Units




Figure 1: A neural network with two input units and one output unit.

Gradient Descent for a Network with Two Input Units

This network can learn up to two associations. Each association consists of an input, which is a pair of values x_1 and x_2 , and each corresponding output is a single value y .

Given one input (x_1, x_2) , the output y is found by multiplying each input value by its corresponding weight and then summing the resultant products:

$$y = w_1x_1 + w_2x_2. \quad (12)$$

Gradient Descent for a Network with Two Input Units

This can be written succinctly if we represent the weights as a vector, written in bold typeface:

$$\mathbf{w} = (w_1, w_2). \quad (13)$$

Similarly, each pair of input values can be represented as a vector, again in bold typeface:

$$\mathbf{x} = (x_1, x_2). \quad (14)$$

The state y for an input \mathbf{x} is found from the *dot* or *inner* product,

$$y = \mathbf{w} \cdot \mathbf{x}, \quad (15)$$

which is defined by Equation 12.

Notice that scalar variables are in italics, whereas vectors are in bold typeface.

Gradient Descent for a Network with Two Input Units

We use subscripts to denote each association

$$\begin{aligned}y_1 &= \mathbf{w} \cdot \mathbf{x}_1, \\ y_2 &= \mathbf{w} \cdot \mathbf{x}_2.\end{aligned}\tag{16}$$

We will write the problem out in full

$$\begin{aligned}y_1 &= w_1x_{11} + w_2x_{21}, \\ y_2 &= w_1x_{12} + w_2x_{22}.\end{aligned}\tag{17}$$

We can recognise this as two simultaneous equations with two unknowns (w_1 and w_2), so we know that a solution for w_1 and w_2 usually exists. We could find this solution manually, but because we know that the problems we encounter later will become unrealistic for manual methods, we will stick to using gradient descent.

Gradient Descent for a Network with Two Input Units

To use gradient descent, we first need to write down an error function like Equation 7. The error function for the first association is

$$E_1 = \frac{1}{2} (\mathbf{w} \cdot \mathbf{x}_1 - y_1)^2, \quad (18)$$

and for the second association it is

$$E_2 = \frac{1}{2} (\mathbf{w} \cdot \mathbf{x}_2 - y_2)^2. \quad (19)$$

The error function for the set of two associations is the sum

$$E = E_1 + E_2 \quad (20)$$

$$= \frac{1}{2} [(\mathbf{w} \cdot \mathbf{x}_1 - y_1)^2 + (\mathbf{w} \cdot \mathbf{x}_2 - y_2)^2], \quad (21)$$

which can be written succinctly using the summation convention as

$$E = \frac{1}{2} \sum_{t=1}^2 (\mathbf{w} \cdot \mathbf{x}_t - y_t)^2. \quad (22)$$

Gradient Descent for a Network with Two Input Units




Figure 2: The error surface is obtained by evaluating Equation 22 over a range of values for w_1 and w_2 . Given an initial weight vector $\mathbf{w} = (-0.8, -1.6)$, the direction of steepest descent is $-\nabla_{\mathbf{w}} E$ (shown by an arrow on the ground plane). The white dots depict the evolution of weights during learning using Equation 30.33

Gradient Descent for a Network with Two Input Units

Using the *chain rule*, the gradient of the error function with respect to w_1 for the t th association ($t = 1$ or 2) is

$$\frac{\partial E_t}{\partial w_1} = \frac{\partial E_t}{\partial y_t} \frac{\partial y_t}{\partial w_1}, \quad (23)$$

where $\partial E_t / \partial y_t = (\mathbf{w} \cdot \mathbf{x}_t - y_t)$, $y_t = \mathbf{w} \cdot \mathbf{x}_t$, and $\partial y_t / \partial w_1 = x_{1t}$, so

$$\frac{\partial E_t}{\partial w_1} = (\mathbf{w} \cdot \mathbf{x}_t - y_t) x_{1t}. \quad (24)$$

Given that the delta term for the t th association is

$$\delta_t = (\mathbf{w} \cdot \mathbf{x}_t - y_t), \quad (25)$$

we then have

$$\frac{\partial E_t}{\partial w_1} = \delta_t x_{1t}. \quad (26)$$

Gradient Descent for a Network with Two Input Units

When considered over both associations, the gradient of the error function with respect to w_1 is

$$\frac{\partial E}{\partial w_1} = \sum_{t=1}^2 \delta_t x_{1t}. \quad (27)$$

Similarly, the gradient with respect to w_2 is

$$\frac{\partial E}{\partial w_2} = \sum_{t=1}^2 \delta_t x_{2t}. \quad (28)$$

Gradient Descent for a Network with Two Input Units

The direction of steepest ascent is a vector on the ground plane that points in the direction to go in order to increase the value of E as quickly as possible. This direction is represented by the *nabla* symbol (∇), which is a vector of scalar gradients:

$$\nabla_{\mathbf{w}} E = \left(\frac{\partial E}{\partial w_1}, \frac{\partial E}{\partial w_2} \right), \quad (29)$$

where the subscript \mathbf{w} indicates a derivative with respect to \mathbf{w} .

If the direction of steepest ascent is ∇E then the direction of steepest descent is $-\nabla E$, as shown in Figure 2. Accordingly, if the current value of the weight vector is \mathbf{w}_{old} , then

$$\mathbf{w}_{\text{new}} = \mathbf{w}_{\text{old}} - \epsilon \nabla E. \quad (30)$$

Gradient Descent for a Network with Two Input Units

```
initialise weights  $\mathbf{w}$  to random values; set learning to true
while learning do
  set recorder of weight change vectors  $\Delta\mathbf{w}$  to zero
  foreach association from  $t = 1$  to 2 do
    set input unit states  $\mathbf{x}$  to  $t$ th training vector  $\mathbf{x}_t$ 
    get state of output unit  $y_t = \mathbf{w} \cdot \mathbf{x}_t$ 
    get delta term  $\delta_t = y_t - y_t$ 
    get weight gradient for  $t$ th input vector  $\nabla E_t = \delta_t \mathbf{x}_t$ 
    get change in weights for  $t$ th input vector  $\Delta\mathbf{w}_t = -\epsilon \nabla E_t$ 
    accumulate weight changes in  $\Delta\mathbf{w} \leftarrow \Delta\mathbf{w} + \Delta\mathbf{w}_t$ 
  end
  update weights  $\mathbf{w} \leftarrow \mathbf{w} + \Delta\mathbf{w}$ 
  if gradient  $|\nabla E| \approx 0$  then
    set learning to false
  end
end
```

A Network with Four Input and Output Units




Figure 3: (a) A neural network with four input units and four output units (i.e. a 4-4 network) can be viewed as four neural networks like the one in (b), where each of the four neural networks has the same four input units but a different output unit (i.e. four 4-1 networks).


Learning Photographs



Learning photographs using a network with the same type of architecture as in Figure 3a. Each photograph \mathbf{x}_t consists of 50×50 pixels. A linear network was used, with an array of 50×50 input units and an array of 50×50 output units. The network was trained to associate each of two training vectors (a and d) with itself.

For example, adding noise to (a) yielded (b), and when (b) was used as input to the network, the output was (c), showing that the network's memory is *content addressable*.

A Backprop Neural Network




Unit states are indexed by layer (x = input, z = hidden, y = output) and by association number t .

A key property of backprop networks is that they can compute nonlinear functions, because most units have nonlinear activation functions.

BUT this also means that the error surface is not convex, so it has *local minima*.

Unit Activation Function




Sigmoidal activation function

$$y = f(u) = (1 + e^{-u})^{-1}, \quad (31)$$

where u is the total input to a unit. The derivative is

$$\frac{dy}{du} = y(1 - y). \quad (32)$$

A Backprop Neural Network



We know how to find the delta term of each output unit. The delta term of the j th hidden unit is its derivative times a weighted sum of delta terms in the output layer

$$\delta_{jt} = \frac{dy_{jt}}{du_{jt}} \sum_{k=1}^K \delta_{kt} w_{jk}, \quad (33)$$

Given the delta term for one association of a unit, the change in an 'incoming' weight for that association is

$$\Delta w_{ij}(t) = -\epsilon \delta_{jt} x_{it}. \quad (34)$$

And, when considered over all T associations, the change in one weight is

$$\Delta w_{ij} = -\epsilon \sum_{t=1}^T \delta_{jt} x_{it}. \quad (35)$$

This is a general recipe for updating weights in a backprop network. Once we have the delta term for each unit then we can adjust all the weights between that unit and all the units in the previous layer.

The Backprop Algorithm

Backprop (Short Version)

initialise network weights \mathbf{w} to random values; set learning to true

while *learning* **do**

 set vector of gradients ∇E to zero

foreach *association from $t = 1$ to N* **do**

 set input unit states \mathbf{x}_{it} to t th training vector

 get state of output units \mathbf{y}_{kt}

 get delta term δ_{kt} for each output unit

 use output delta terms to get hidden unit delta terms

 use delta terms to get vector of weight gradients ∇E_t

 accumulate gradient $\nabla E \leftarrow \nabla E + \nabla E_t$

end

 get weight change $\Delta \mathbf{w} = -\epsilon \nabla E$

 update weights $\mathbf{w} \leftarrow \mathbf{w} + \Delta \mathbf{w}$

if $|\nabla E| \approx 0$ **then**

 set learning to false

end

end

Global and Local Minima




Figure 4: Schematic diagram of local and global minima in the error function $E(\mathbf{w})$ for a network with just two weights.

Global and Local Minima

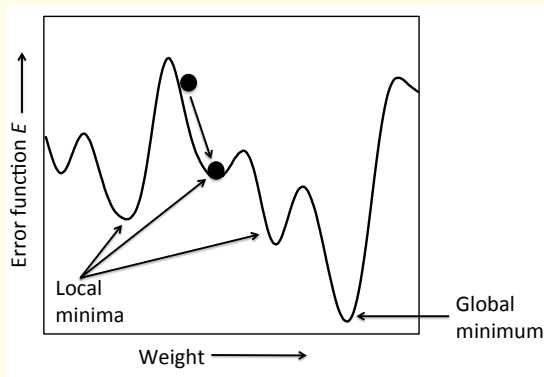




Figure 5: Local and global minima in a cross-section of the error function $E(\mathbf{w})$. At a given initial weight, $E(\mathbf{w})$ may be high, as represented by the black disc. Gradient-based methods always head downhill, but because they can only move downwards, the final weight often corresponds to a local minimum.


Over-fitting, under-fitting, Goldilocks-fitting



Underfitted




Good Fit/Robust



Overfitted

Preventing Over-fitting with Early Stopping



While learning a *training set*, the error on a separate *validation set* is monitored to gauge *generalisation* performance. Training is stopped when the validation error stops decreasing. Other methods include *regularisation*.

The End

Thank you.